Real-Time Collaborative Editor Design Documentation

1. Overview, Requirements, and Assumptions

A. Functional Requirements:

- Real-Time Collaboration: Multiple users can edit the same document simultaneously with immediate updates.
- Rich-Text Editing: Support basic and advanced formatting, inline images, and embedded media.
- Conflict Resolution: Merge concurrent edits using Operational Transformation (OT) or CRDTs.
- Access Control and Sharing: Role-based permissions for document sharing and collaboration.
- Offline and Resynchronization: Allow offline edits with automatic synchronization upon reconnection.
- Version Control: Maintain complete document version history with rollback support.

B. Nonfunctional Requirements:

- Low Latency: Edits propagate within a few hundred milliseconds across clients.
- Scalability: Support millions of documents and thousands of concurrent collaborative sessions.
- High Availability: Distributed, fault-tolerant architecture with minimal disruption on failures.
- Global Distribution: Regional data centers deliver low-latency service worldwide.
- Security: TLS encryption for communications and strong authentication and access control.

C. Assumptions:

- Billions of registered users; only a fraction concurrently edit documents.
- Documents are represented as sequences of operations that transform a base state.
- The system leverages a microservices architecture with a dedicated synchronization service.

2. High-Level Architecture and Component Responsibilities

A. Client Tier:

- Editor Clients: Web, mobile, and desktop apps use HTTPS for API calls and secure WebSocket (WSS) for real-time updates.
- Local Edit Buffer: Captures local edits and displays real-time changes.

B. Global Access:

- Global DNS and Regional Load Balancers: Route users to the nearest data center for low latency.

C. API Gateway and Authentication:

- API Gateway: Authenticates incoming requests, applies rate limiting, and routes them to the collaboration services.

- Authentication Service: Issues tokens (e.g., JWT) to manage secure sessions.

D. Real-Time Collaboration Service:

- Collab Server: Maintains the real-time session for each document, receives client operations, and applies conflict resolution (using OT or CRDT).

- OT/CRDT Engine: Transforms concurrent operations to ensure all clients converge on the same document state.

- Presence and Cursor Tracking: Broadcasts collaborators' cursors and selections for real-time awareness.

E. Persistence and Storage:

- Document Store: A distributed NoSQL or relational database that persists document states and version histories.

- Operation Log and Snapshot Service: Periodically checkpoints document states and stores operation logs for recovery.

F. Messaging and Background Processing:

- Message Broker: Decouples real-time operations from background processing (e.g., snapshot creation) using pub/sub mechanisms.

G. Analytics and Audit:

- Logging and Monitoring: Tracks performance, concurrent edits, and sync latency; supports troubleshooting and auditing.

3. Detailed Workflow

A. Connection and Document Loading:

1. A user logs in via HTTPS; the API Gateway authenticates and issues a token.
2. The client requests a specific document; the Collab Server retrieves the latest document state (via snapshots and operation logs) from the Document Store.
3. A secure WebSocket connection is established between the client and the Collab Server.

B. Real-Time Editing and Synchronization:

1. Local edits (operations) are captured by the client and sent over the WebSocket to the Collab Server.
2. The Collab Server applies conflict resolution (OT/CRDT) to merge concurrent edits.
3. Transformed operations are broadcast to all connected clients for real-time update.
4. The document state is updated in-memory and periodically checkpointed to the Document Store.

C. Offline Support and Resynchronization:

- The client buffers local changes during disconnection and synchronizes them when reconnected, with the Collab Server handling potential conflicts.

D. Version Control and Audit:

- All operations are logged, enabling users to review change history and revert to previous versions if necessary.

4. Scalability, Fault Tolerance, and Global Distribution

A. Horizontal Scalability:

- Collab Servers are partitioned by document ID, allowing each server to manage a subset of active sessions.
- Stateless API Gateways and Authentication Services are scaled dynamically with autoscaling groups.
- The Document Store and operation logs are sharded to manage a high volume of documents and operations.

B. Fault Tolerance and High Availability:

- Document state and operation logs are replicated across multiple nodes and regions.
- In case of a Collab Server failure, client connections are rerouted to another instance with minimal disruption.

C. Global Distribution:

- Users are served from the nearest regional data center via DNS routing and global load balancing.
- Multi-region deployment minimizes latency and supports disaster recovery.

5. Protocols and Security

A. Communication Protocols:

- Client-to-Server: HTTPS for API calls and secure WebSocket (WSS) for real-time operations.
- Interservice Communication: gRPC or REST over secured TCP with mutual TLS for low-latency data exchange.

B. Security and Access Control:

- All data in transit is encrypted (TLS); sensitive data at rest is also encrypted.
- Role-based access control (RBAC) restricts document editing and sharing.

- Audit logs record all operations for compliance and troubleshooting.

6. Final Thoughts

This design for a real-time collaborative editor (similar to Google Docs) presents a robust, scalable architecture that supports real-time editing, conflict resolution, and version control. Key highlights include:

- Real-time synchronization using OT/CRDT techniques that enable concurrent editing with low latency.
- A microservices-based approach that separates collaboration, persistence, and authentication layers for high availability and scalability.
- Global distribution with regional load balancing to reduce latency and improve performance.
- Built-in offline support, comprehensive version history, and secure access controls.

This framework provides a solid conceptual foundation for building a collaborative editing platform that can scale to millions of concurrent sessions and support rich, real-time user interactions.