

# Scalable Netflix-Like System Design Documentation

## 1. Overview, Requirements, and Assumptions

### A. Functional Requirements:

- Video Streaming on Demand: Provide high-quality, adaptive bitrate streaming across multiple devices.
- Content Delivery & Playback: Enable low-latency, instant playback with features such as pause, rewind, and fast-forward.
- User Account and Profile Management: Support account creation, watch lists, and billing/subscription services.
- Personalization and Recommendation: Tailor content suggestions based on user viewing history and preferences.
- Search and Discovery: Provide advanced search and content categorization.
- Content Management: Ingest, transcode, and manage metadata for thousands of titles.

### B. Nonfunctional Requirements:

- Low Latency and High Throughput: Ensure playback starts within seconds and supports millions of concurrent streams.
- High Availability and Fault Tolerance: Achieve near-100% uptime with multi-region redundancy.
- Global Distribution: Deliver content worldwide via a CDN and regional edge servers.
- Security: Secure content delivery with TLS and protect user data with strong authentication and DRM.

### C. Assumptions and Scale:

- Billions of registered users; hundreds of millions active.
- Peak concurrent streaming sessions in the tens to hundreds of millions.
- Content library of thousands of titles, each with multiple encoded versions.
- Deployment across hundreds of data centers and tens of thousands of servers globally.

## 2. High-Level Architecture and Component Responsibilities

### A. Client Tier:

- Applications on Smart TVs, mobile, and desktop use HTTPS for API calls and protocols like MPEG-DASH/HLS for video streaming.

### B. Global Access and Routing:

- DNS-based routing and a robust CDN deliver content to users from the nearest edge server.
- Regional load balancers distribute incoming requests among data centers.

# Scalable Netflix-Like System Design Documentation

## C. API Gateway and Service Layer:

- API Gateway handles authentication, session management, and routing via HTTPS and secure protocols (gRPC/REST over mTLS).

- Authentication and Billing Services manage user accounts and subscriptions.

## D. Content Ingestion and Transcoding:

- New content is ingested, transcoded into multiple resolutions/bitrates using specialized clusters, and stored in origin servers.

- Metadata and DRM services apply content protection and generate searchable indexes.

## E. Video Delivery and Playback:

- Origin and streaming servers deliver video manifests (MPEG-DASH/HLS) and segments.

- Adaptive bitrate streaming ensures smooth playback based on network conditions.

## F. Personalization, Search, and Analytics:

- Recommendation engines and search services use user data to personalize and facilitate content discovery.

- Data stores and analytics platforms capture viewing metrics and drive optimization.

## G. External Infrastructure and Security:

- A CDN handles global content delivery while data centers provide origin storage and streaming.

- Optional Distributed ID Generators, based on Twitter Snowflake, ensure globally unique identifiers.

- DRM protects content and secure protocols (TLS, mTLS) safeguard all communications.

## 3. Detailed Workflow

### A. User Request and Playback:

- Users log in via HTTPS; the API Gateway authenticates and issues tokens.

- The client selects a video and receives metadata and a streaming manifest.

- Playback is initiated with adaptive streaming from the nearest CDN edge.

### B. Content Ingestion and Transcoding:

- New titles are ingested, transcoded into multiple formats, and stored on origin servers, then pushed to CDN caches.

### C. Personalization and Recommendations:

- User interactions and viewing history feed data into recommendation engines to generate personalized content lists.

# Scalable Netflix-Like System Design Documentation

## D. Analytics and Monitoring:

- Real-time analytics track playback performance, user engagement, and system health for ongoing optimization.

## 4. Scalability, Fault Tolerance, and Global Distribution

### A. Horizontal Scalability:

- Thousands of edge nodes (CDN) serve content; origin and streaming servers scale horizontally across hundreds of data centers.
- API Gateways, content ingestion, and transcoding clusters are auto-scaled based on demand.

### B. Fault Tolerance and Replication:

- Content is replicated across regions with a replication factor of 3 or more.
- Redundancy in API services, streaming servers, and data stores ensures no single point of failure.

### C. Global Distribution:

- DNS-based routing and regional load balancing direct users to the nearest available servers, ensuring low latency.
- Multi-region deployment supports disaster recovery and continuous service during regional outages.

## 5. Protocols and External Infrastructure

### A. Communication Protocols:

- Client-to-Server: HTTPS for API calls; MPEG-DASH/HLS over HTTPS for streaming.
- Interservice: gRPC or REST over secured TCP (with mutual TLS) for low-latency service interactions.

### B. External Infrastructure Components:

- CDN: A global content delivery network caches video segments at edge servers worldwide.
- Distributed ID Generator (Optional): A microservice (using the Twitter Snowflake algorithm) generates globally unique IDs if needed for session or content management.
- DRM: External DRM services protect video content from unauthorized access.

### C. Load Balancing and Autoscaling:

- Global DNS and regional load balancers distribute user requests effectively.
- Autoscaling policies monitor load metrics to dynamically provision additional server resources as

# Scalable Netflix-Like System Design Documentation

needed.

## 6. Final Thoughts

This design for a Netflix-like streaming service embodies the principles necessary to deliver high-quality video to billions of users worldwide. Key elements include:

- A robust, globally distributed infrastructure with CDNs and regional data centers.
- Scalable microservices for API processing, content ingestion, transcoding, and personalized recommendations.
- A data persistence layer that combines durable storage (NoSQL, object stores) with in-memory caching for performance.
- Secure, efficient communication protocols ensuring both data integrity and low latency.

While Netflix's real-world system includes many proprietary optimizations and highly specialized hardware, this conceptual design provides a strong architectural foundation to achieve similar levels of scalability, resilience, and user experience.